

SHL AI Hiring Submission Package

A compact, submission-ready summary for the assessment portal. It answers the form prompts and summarizes the implementation and verification work.

Form answers

Did the solution meet the expectations?	Yes. The app asks clarifying questions, returns SHL recommendations, refines the shortlist, compares assessments, and includes evaluation checks.
Public base URL	Not deployed permanently from this workspace. Use Railway or Render to create a stable URL before final submission.
Cold start delay	N/A until deployed. On free tiers that sleep, expect a short first-request delay. Use an always-on plan to avoid it.
LLM used	Gemini 2.5 Flash via the Gemini API.
AI tools used	OpenAI Codex for implementation, debugging, and testing; Gemini API for response polishing.

Submission note

The public base URL should be filled in separately after deployment. This PDF intentionally avoids hosting instructions.

Approach summary

- Built a stateless FastAPI service. Each /chat request includes the transcript so the backend does not rely on session memory.
- Used the SHL catalog JSON endpoint as the source of truth, cached normalized records, and indexed the catalog for grounded retrieval.
- The planner classifies requests into clarify, recommend, compare, or refuse, then applies catalog-driven filters and comparison logic.
- Prompt templates keep the assistant concise and recruiter-like. Gemini is only used to polish the final reply, not to invent recommendations.
- Evaluation uses pytest scenarios plus a smoke script covering clarification, recommendation, refinement, comparison, and guardrail behavior.
- What did not work initially: Gemini sometimes returned JSON wrapped in code fences. I normalized the reply text so the UI receives clean recruiter language.
- Measured improvement by running the full test suite, smoke checks, and rendering the PDF output to verify layout quality.

Evaluation notes

- The regression suite passed end to end after the Gemini reply normalization fix.
- A smoke run confirmed clarification, recommendation, refinement, comparison, and guardrail behavior.
- The PDF output was rendered to images and visually checked for layout issues.
- The UI serves the chat experience from the root route and stays stateless across turns.
- The final implementation remains grounded in the catalog JSON and does not invent assessments outside the source data.

If you need the deployment steps separately, they can be provided outside this PDF so the submission file stays focused on the solution summary.